

INTRODUCTION TO MPI

Four interlocking puzzle pieces are arranged in a cluster on the right side of the slide. The pieces are colored red, yellow, blue, and green. They are slightly offset from each other, creating a sense of depth.

Dany Vohl

Astronomy Data and Computing Services,
Swinburne

TABLE OF CONTENT

- Introduction
 - Notions of parallel computing
 - *Shared Memory versus Distributed Memory*
 - Message Passing Interface
- Working examples with mpi4py
 - Point-to-Point Communication
 - *Sum two vectors*
 - Collective Communication Patterns
 - *Computing Pi*

DISCLAIMER

- We chose Python over C/C++ for simplicity and to avoid the hassle of compilations and Linking.
- The code in some cases is not optimized or follow best practices.
 - *The important point is to illustrate the concept.*
- Most of these tools are actually designed and developed for C/C++.
 - *Their python version might not be complete.*




INTRODUCTION

*Shared Memory VS
Distributed Memory*

SHARED MEMORY VERSUS DISTRIBUTED MEMORY


SHARED MEMORY



Modified From: Hanjun Kim, Princeton University

SHARED MEMORY VERSUS DISTRIBUTED MEMORY


SHARED MEMORY



Modified From: Hanjun Kim, Princeton University

SHARED MEMORY VERSUS DISTRIBUTED MEMORY


SHARED MEMORY



Modified From: Hanjun Kim, Princeton University

SHARED MEMORY VERSUS DISTRIBUTED MEMORY

SHARED MEMORY




Modified From: Hanjun Kim, Princeton University

SHARED MEMORY VERSUS DISTRIBUTED MEMORY


SHARED MEMORY

- Splitting the problem among many people in shared memory
 - Pros - Easy to use
 - Cons - Limited Scalability, High coherence overhead




SHARED MEMORY VERSUS DISTRIBUTED MEMORY

SHARED MEMORY



SHARED MEMORY VERSUS DISTRIBUTED MEMORY

DISTRIBUTED MEMORY




Modified From: Hanjun Kim, Princeton University

SHARED MEMORY VERSUS DISTRIBUTED MEMORY

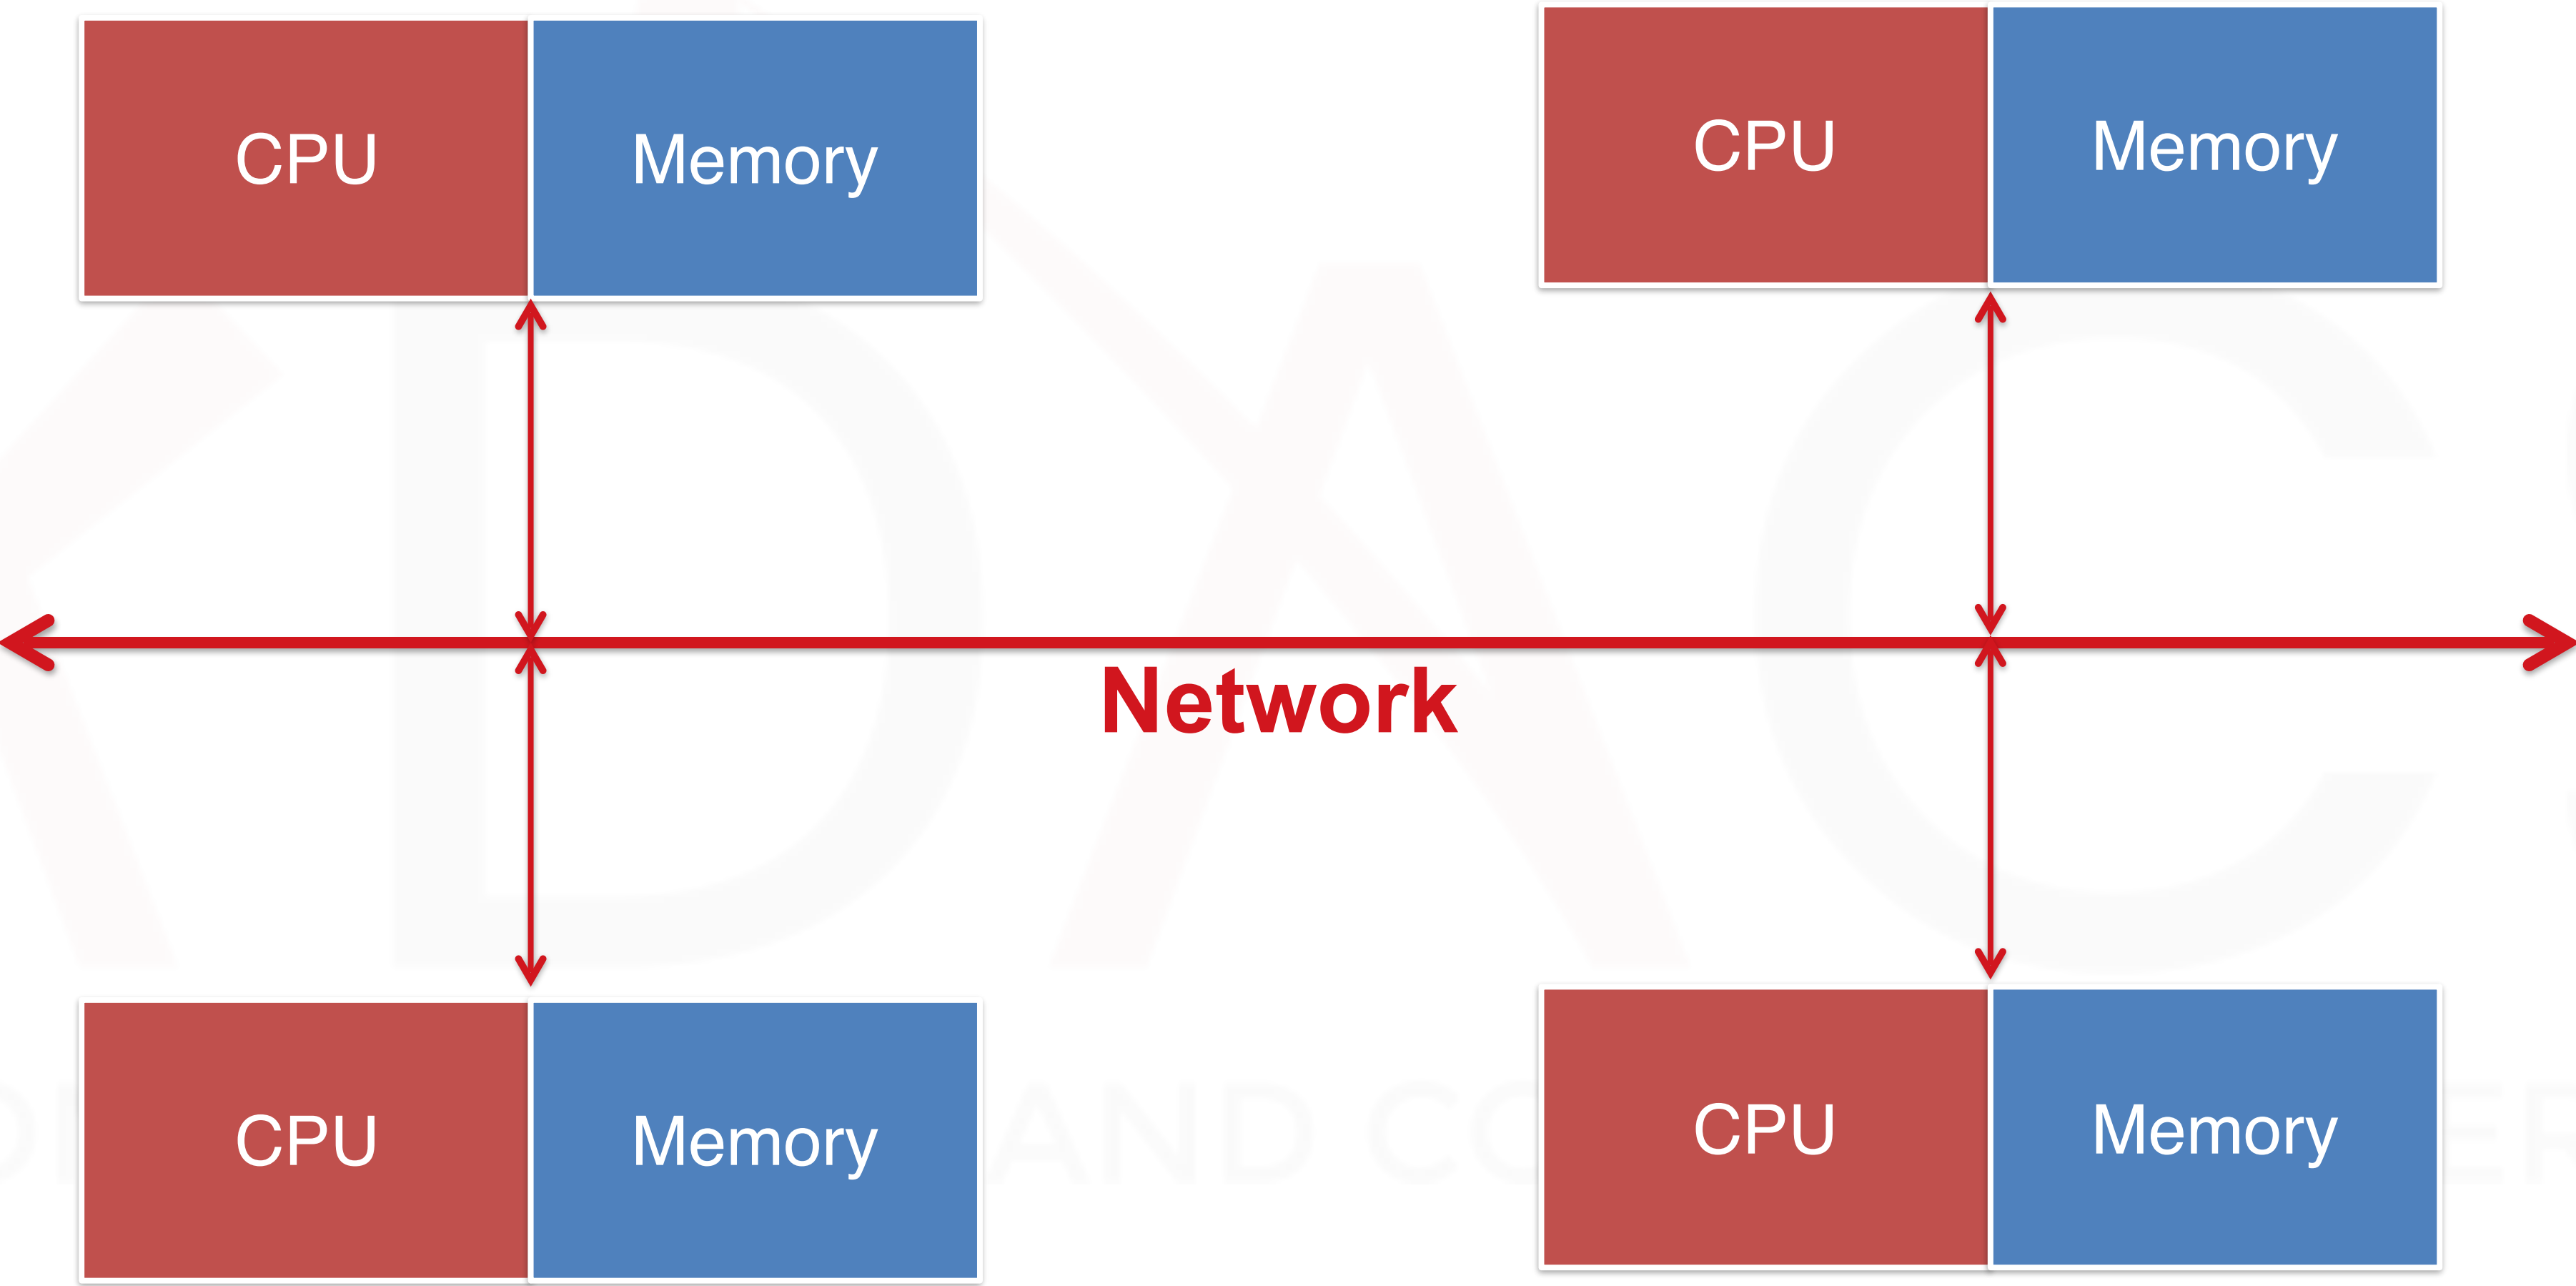
DISTRIBUTED MEMORY

- Splitting the problem among many people distributed memory
 - *Scalable seats (Scalable resources)*
 - *Less contention from private memory spaces*




SHARED MEMORY VERSUS DISTRIBUTED MEMORY

DISTRIBUTED MEMORY



SHARED MEMORY VERSUS DISTRIBUTED MEMORY

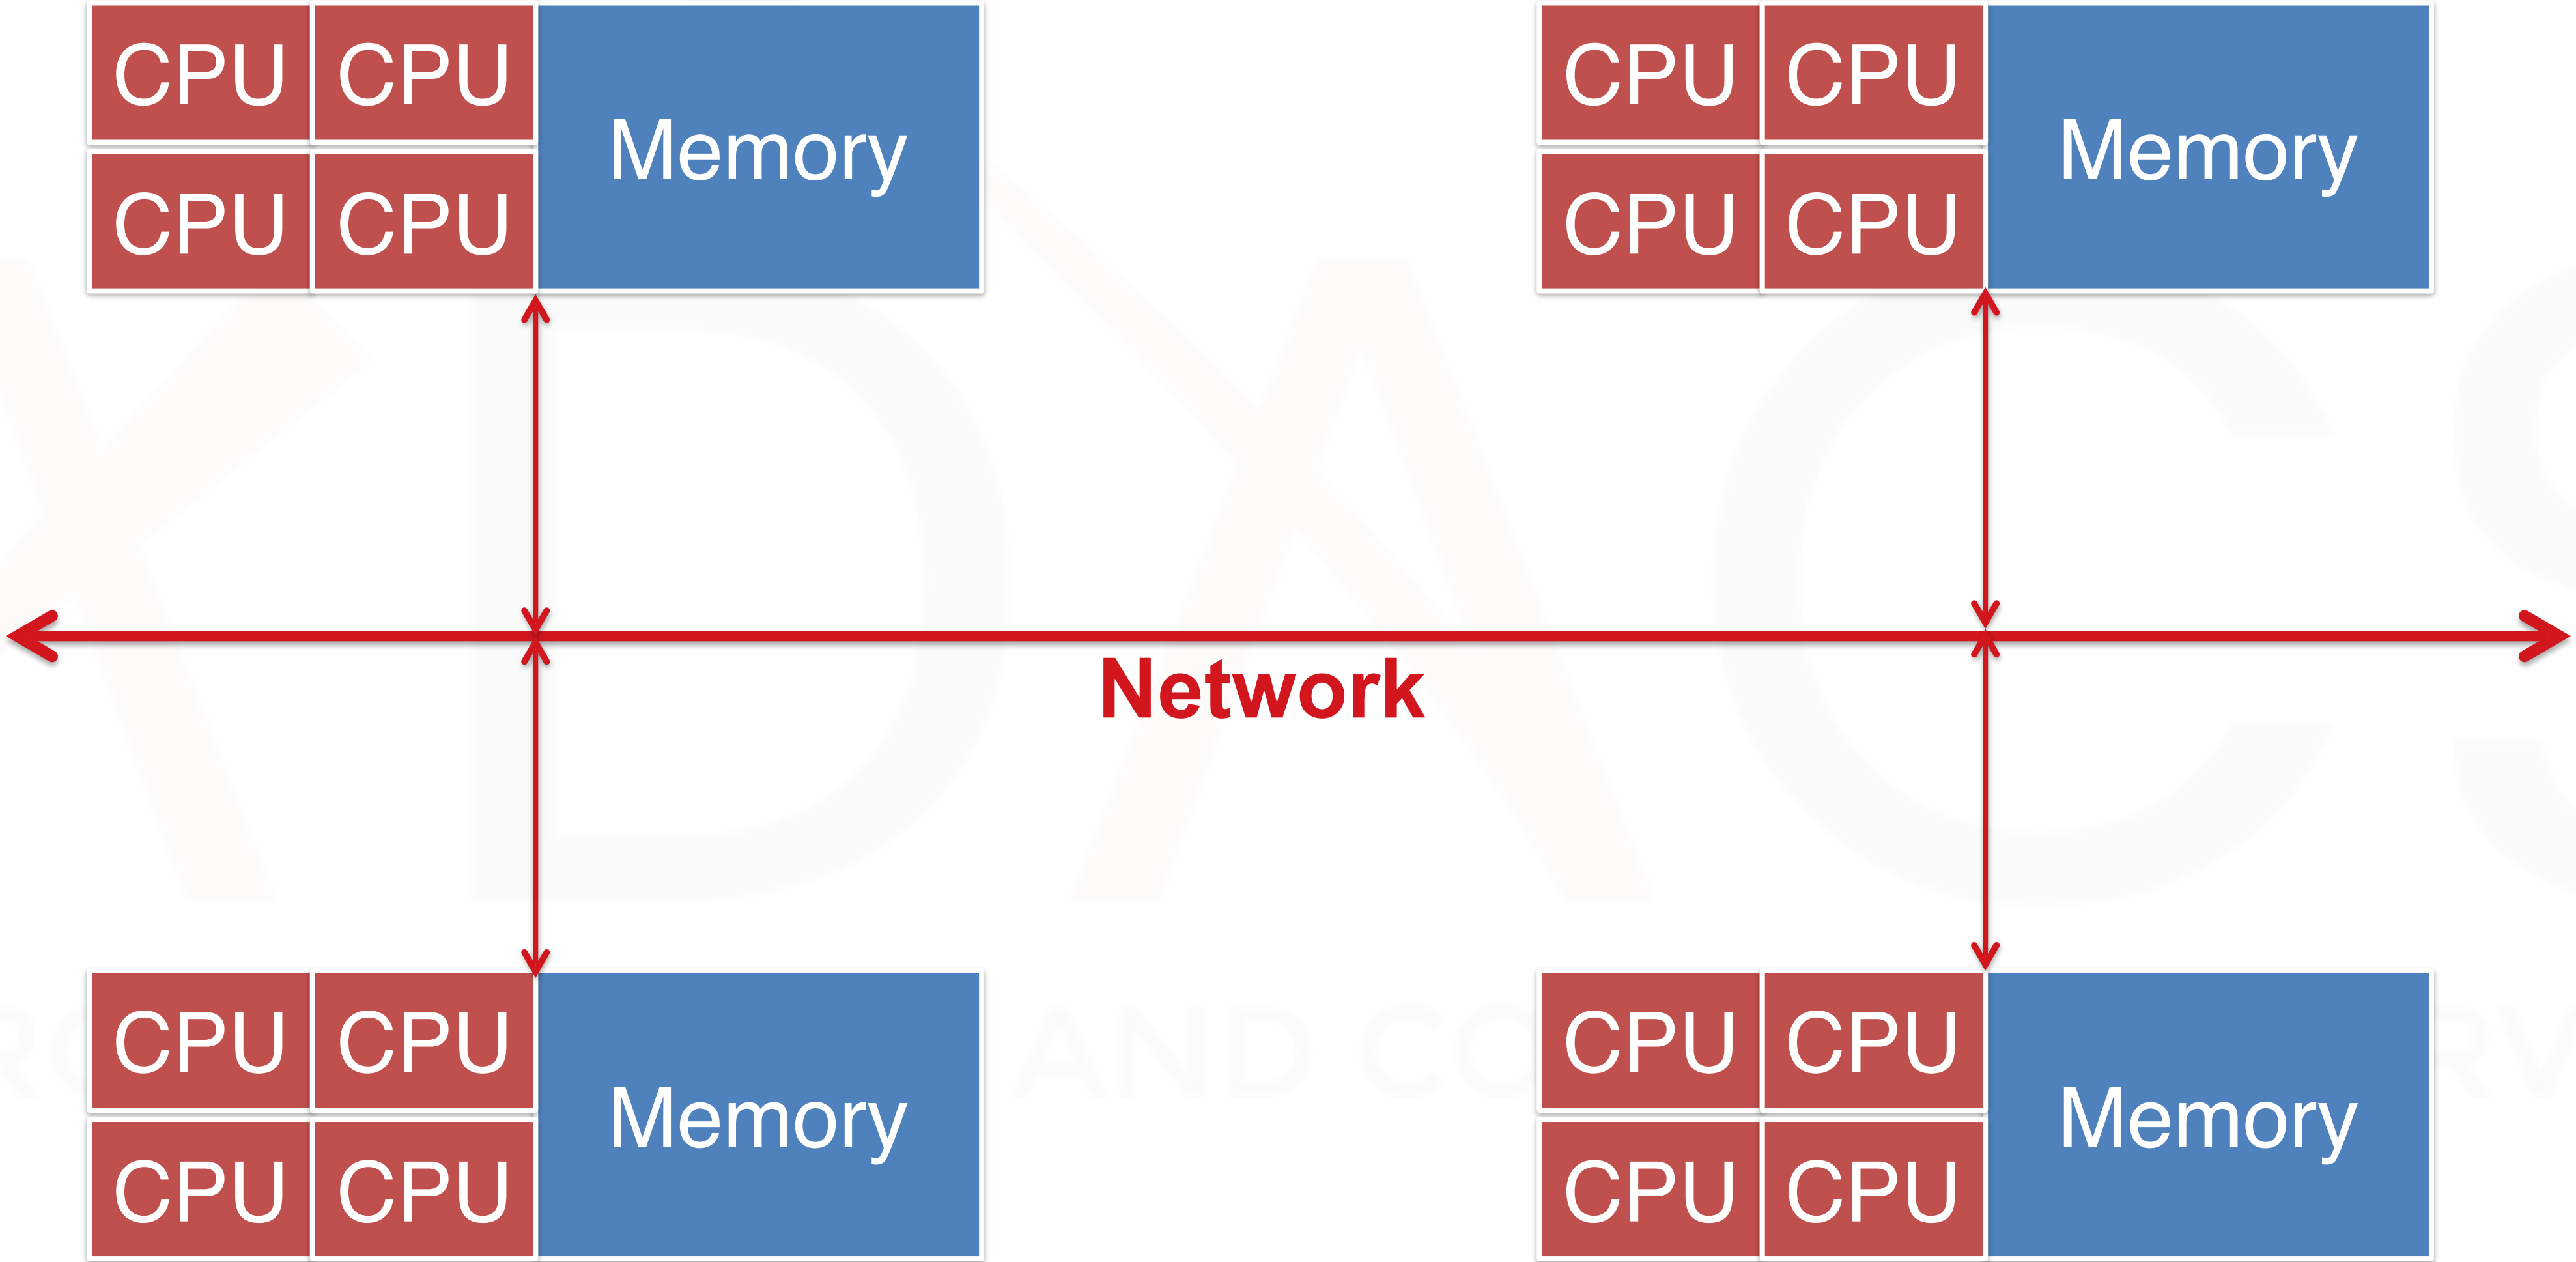
MIXED VERSION



Modified From: Hanjun Kim, Princeton University

SHARED MEMORY VERSUS DISTRIBUTED MEMORY

MIXED VERSION





INTRODUCTION

Message Passing Interface

ASTRONOMY DATA AND COMPUTING SERVICES

WHAT IS MPI, AND WHY USE IT?

*"The Message Passing Interface Standard (MPI) is a message passing library standard based on the consensus of the MPI Forum, which has over 40 participating organizations, including vendors, researchers, software library developers, and users. **The goal of [MPI] is to establish a portable, efficient, and flexible standard for message passing that will be widely used for writing message passing programs.** (...) The advantages of developing message passing software using MPI closely match the design goals of portability, efficiency, and flexibility. **MPI is not an IEEE or ISO standard, but has in fact, become the "industry standard" for writing message passing programs on HPC platforms.**"*

– <https://computing.llnl.gov/tutorials/mpi/>





WHAT IS MPI, AND WHY USE IT?

- De facto Standard Framework for Distributed computing.
- Simple communication model between processes in a program.
- Multiple implementations; highly efficient for different platforms.
- Well established community (since 1994).



ADACS

ASTRONOMY DATA AND COMPUTING SERVICES



MESSAGE PASSING INTERFACE




MESSAGE PASSING INTERFACE



MPI_COMM_WORLD




MESSAGE PASSING INTERFACE




MESSAGE PASSING INTERFACE

Point-to-Point Communication



MESSAGE PASSING INTERFACE

Collective Communication





WORKING EXAMPLES WITH MPI4PY

Message Passing Interface

ASTRONOMY DATA AND COMPUTING SERVICES

MPI FOR PYTHON

- Python package : mpi4py
- mpi4py supports:
 - (1) Pickle-based communication of generic Python object
 - *For convenience*
 - (2) Direct array data communication of buffer-provider objects
 - *Faster option (near C-speed)*
 - (e.g., NumPy arrays).

MPI FOR PYTHON

- Important Class "**Comm**"

1. Communication for Generic python objects

- Use "lower case" methods: **send()**, **receive()**

2. Communication for buffer-provider objects (e.g numpy arrays)

- Use "upper case" methods: **Send()**, **Receive()**

MPI FOR PYTHON

- Importing library:

```
from mpi4py import MPI
```

- Getting important information:

```
comm = MPI.COMM_WORLD
```

```
rank = MPI.COMM_WORLD.Get_rank()
```

```
size = MPI.COMM_WORLD.Get_size()
```

```
name = MPI.Get_processor_name()
```

```
from mpi4py import MPI
comm=MPI.COMM_WORLD
print ("Hello! I'm rank %d out of %d processes running in
total ..."%(comm.rank,comm.size))
comm.Barrier()
```

```
>> mpirun -np 4 python hello-world.py
```

```
from mpi4py import MPI
comm=MPI.COMM_WORLD
print ("Hello! I'm rank %d out of %d processes running in
total ... running on %s"%(comm.rank, comm.size,
MPI.Get_processor_name()))
comm.Barrier()
```

```
>> qsub -q sstar -I -l walltime=0:5:0,nodes=2,ppn=14,
    mem=200MB
```

```
>> mpirun -np 28 python hello-world-2.py
```




WORKING EXAMPLES WITH MPI4PY

*Message Passing Interface
(Point-to-Point Communication)*

MPI FOR PYTHON SENDING AND RECEIVING DATA (POINT-TO-POINT COMM.)

- Python objects (pickle under the hood)
 - *An object to be sent is passed as a parameter to the communication call, and the received object is simply the return value.*

ADACS

ASTRONOMY DATA AND COMPUTING SERVICES

MPI FOR PYTHON SENDING AND RECEIVING DATA (POINT-TO-POINT COMM.)

- Python objects (pickle under the hood)

```
from mpi4py import MPI
```

```
comm = MPI.COMM_WORLD
```

```
rank = comm.Get_rank()
```

```
if rank == 0:
```

```
    data = {'a': 7, 'b': 3.14}
```

```
    comm.send(data, dest=1, tag=11)
```

```
elif rank == 1:
```

```
    data = comm.recv(source=0, tag=11)
```

MPI FOR PYTHON SENDING AND RECEIVING DATA (POINT-TO-POINT COMM.)

- Python objects with non-blocking communication
 - The ***isend()*** and ***irecv()*** methods return *Request* instances
 - Completion of these methods is managed via the ***test()*** and ***wait()*** methods of the *Request* class.

MPI FOR PYTHON SENDING AND RECEIVING DATA (POINT-TO-POINT COMM.)

- Python objects with non-blocking communication

```
from mpi4py import MPI
```

```
comm = MPI.COMM_WORLD
```

```
rank = comm.Get_rank()
```

```
if rank == 0:
```

```
    data = {'a': 7, 'b': 3.14}
```

```
    req = comm.isend(data, dest=1, tag=11)
```

```
    req.wait()
```

```
elif rank == 1:
```

```
    req = comm.irecv(source=0, tag=11)
```

```
    data = req.wait()
```


MPI FOR PYTHON SENDING AND RECEIVING DATA (POINT-TO-POINT COMM.)

- Buffer-provider objects (NumPy arrays; the fast way)
 - *Buffer arguments to these calls must be explicitly specified*
 - *Use a 2/3-list/tuple*
 - e.g. [data, MPI.DOUBLE]
 - uses the byte-size of data and the extent of the MPI datatype to define the count
 - or [data, count, MPI.DOUBLE]

MPI FOR PYTHON SENDING AND RECEIVING DATA (POINT-TO-POINT COMM.)

- Buffer-provider objects (NumPy arrays; the fast way)

```
from mpi4py import MPI
import numpy

comm = MPI.COMM_WORLD
rank = comm.Get_rank()

# passing MPI datatypes explicitly
if rank == 0:
    data = numpy.arange(1000, dtype='i')
    comm.Send([data, MPI.INT], dest=1, tag=77)
elif rank == 1:
    data = numpy.empty(1000, dtype='i')
    comm.Recv([data, MPI.INT], source=0, tag=77)

(...)
```

MPI FOR PYTHON SENDING AND RECEIVING DATA (POINT-TO-POINT COMM.)

- Buffer-provider objects (NumPy arrays; the fast way)

(...)

automatic MPI datatype discovery

if rank == 0:

 data = numpy.arange(100, dtype=numpy.float64)

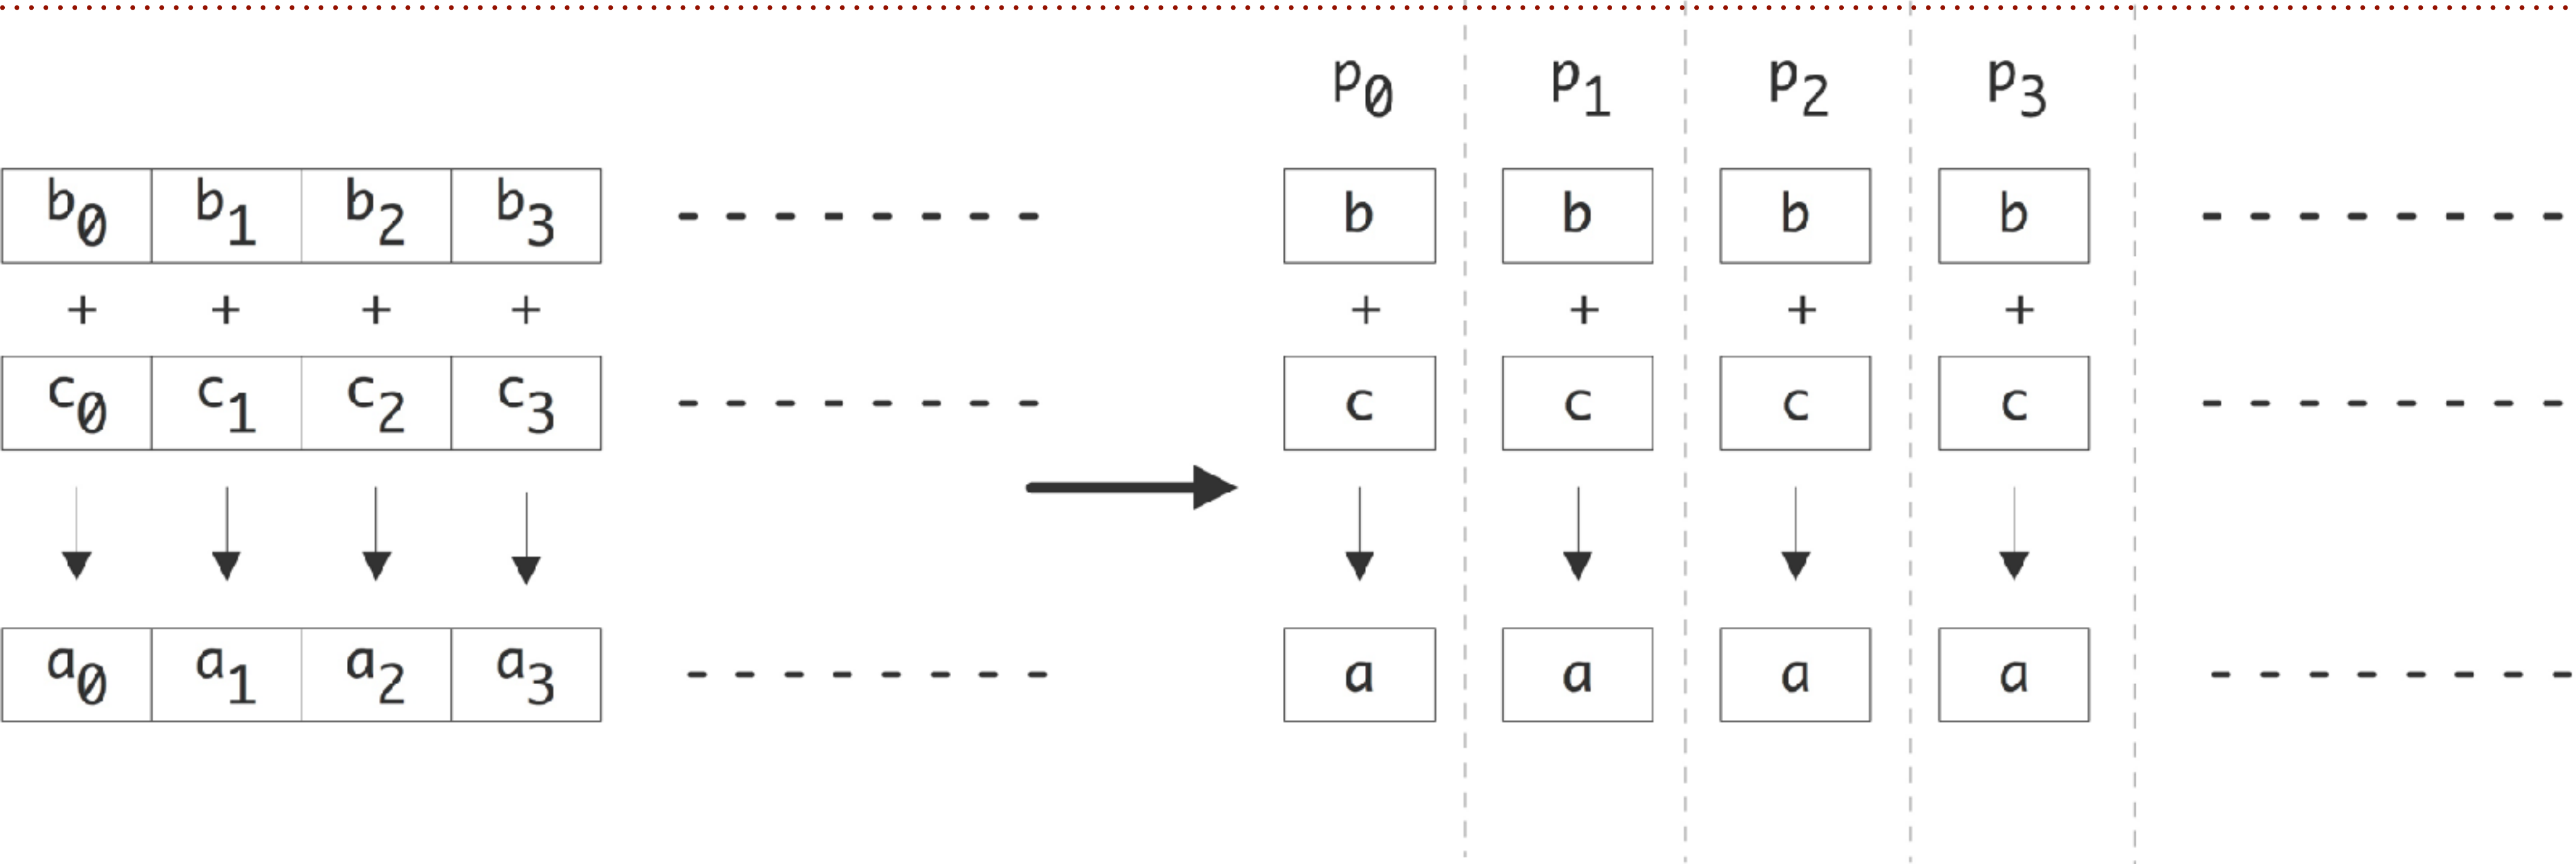
 comm.Send(data, dest=1, tag=13)

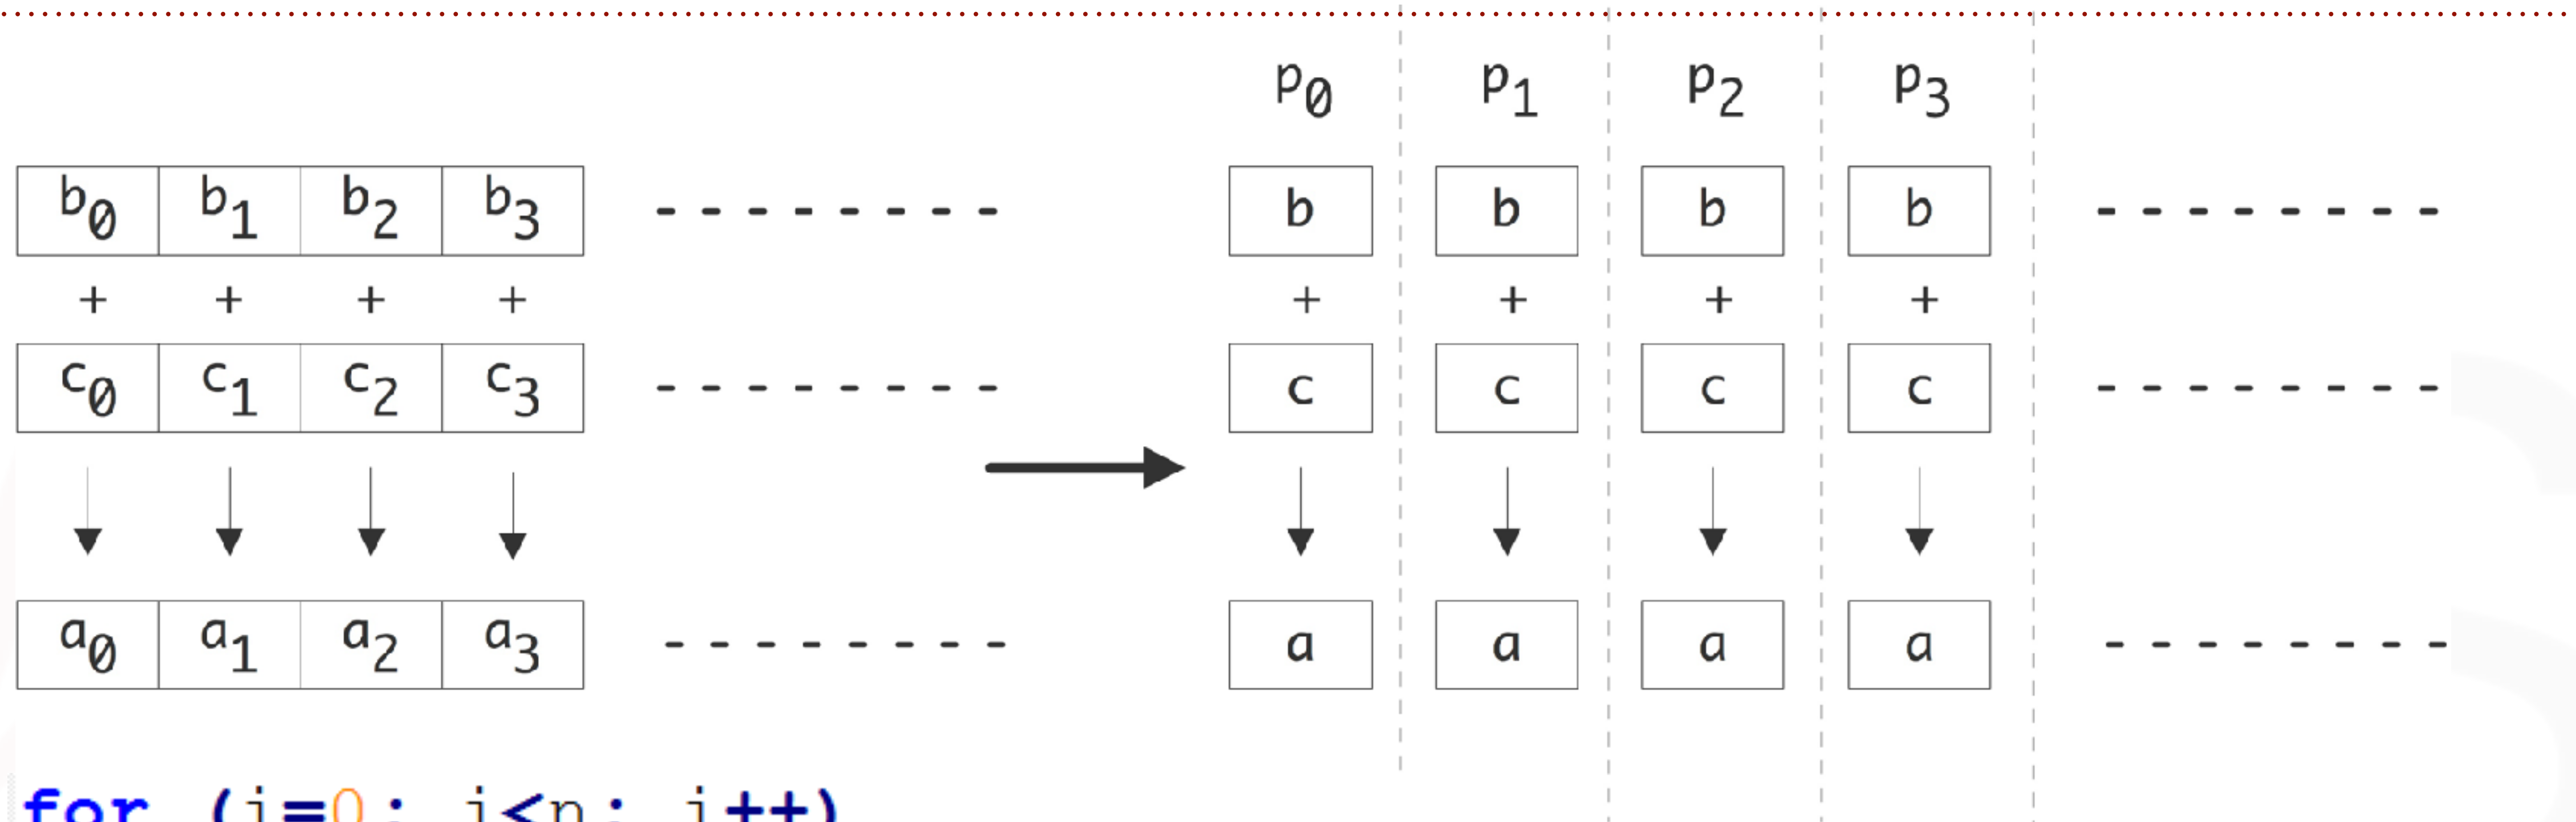
elif rank == 1:

 data = numpy.empty(100, dtype=numpy.float64)

 comm.Recv(data, source=0, tag=13)

ASTRONOMY DATA AND COMPUTING SERVICES






```
for (i=0; i<n; i++)  
    a[i] = b[i] + c[i];
```




```
for (i=my_low; i<my_high; i++)  
    a[i] = b[i] + c[i];
```


MPI FOR PYTHON

SUM TWO VECTORS



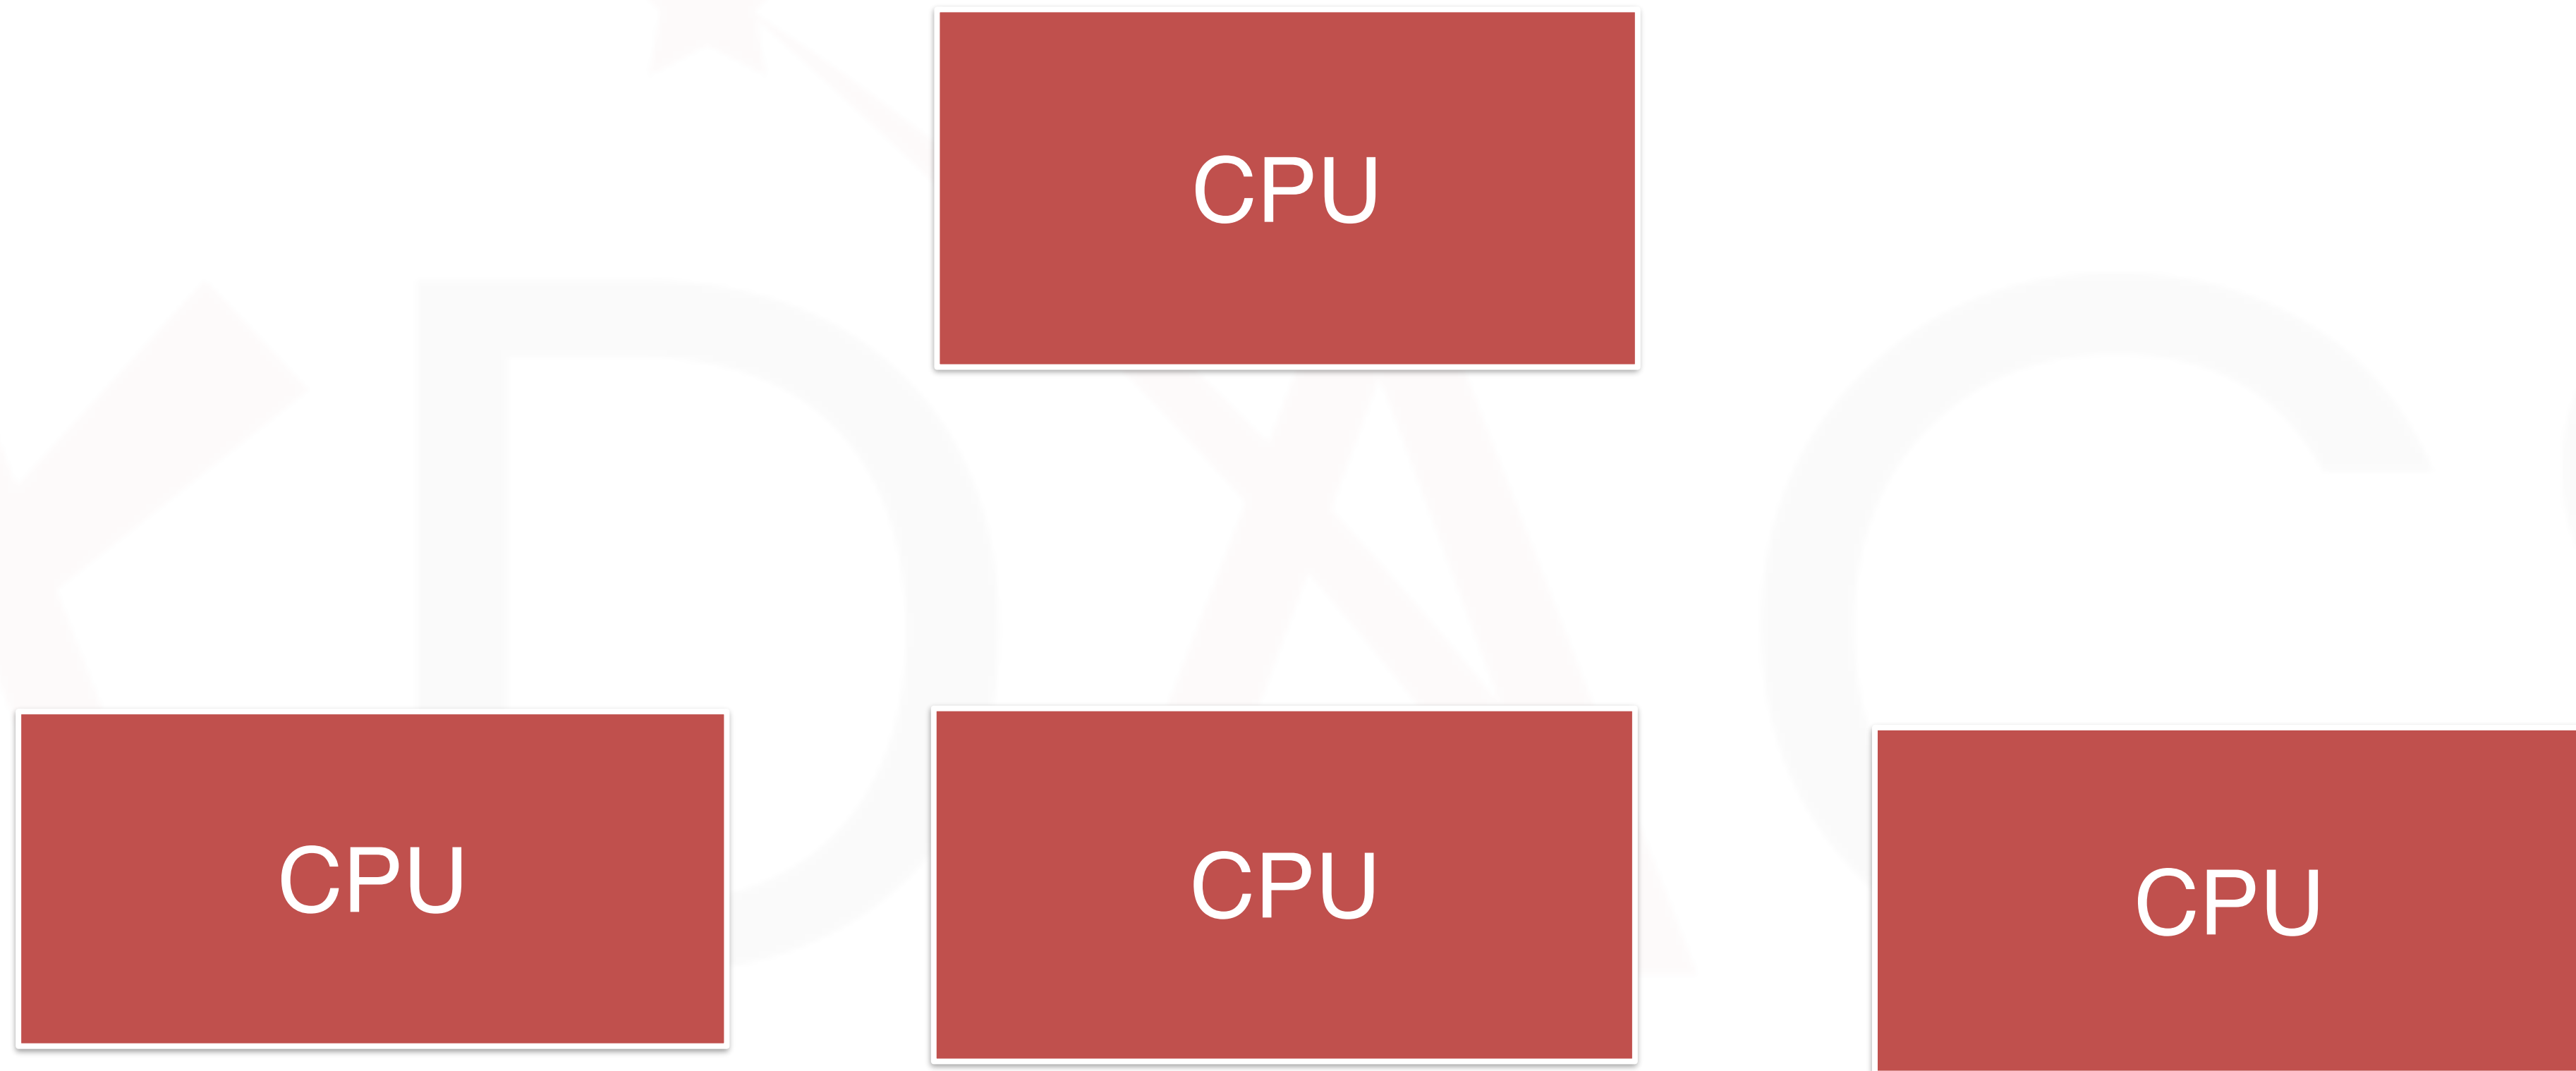
Client-Server Communication




Client-Server Communication

MPI FOR PYTHON

SUM TWO VECTORS



Client-Server Communication



Client-Server Communication

```
from mpi4py import MPI
import numpy

comm = MPI.COMM_WORLD
rank = comm.Get_rank()
size = comm.Get_size()
Count = 999
PCount = Count / (size - 1)

(...)
```



```
(...)  
# pass explicit MPI datatypes  
if rank == 0:  
    Adata = numpy.arange(0, Count, 1, dtype='i')  
    Bdata = numpy.arange(Count, 0, -1, dtype='i')  
    Cdata = numpy.empty(Count, dtype='i')  
    for p in range(1, size):  
        Start = PCount * (p - 1)  
        End = PCount * p  
        print ("%d:[%d to %d]" % (p, Start, End))  
  
        Dim = [Start, End]  
        comm.send(Dim, dest=p, tag=1)  
        comm.Send(Adata[Start:End], dest=p, tag=2)  
        comm.Send(Bdata[Start:End], dest=p, tag=2)  
  
    for p in range(1, size):  
        Start = PCount * (p - 1)  
        End = PCount * p  
        comm.Recv(Cdata[Start:End], source=p, tag=3)  
    print (Cdata)  
(...)
```

(...)
else:

```
Dim = comm.recv(source=0, tag=1)
print ("Rank %d: Received[%d to %d]" % (rank, Dim[0], Dim[1]))

Adata = numpy.empty(Dim[1] - Dim[0], dtype='i')
Bdata = numpy.empty(Dim[1] - Dim[0], dtype='i')

comm.Recv(Adata, source=0, tag=2)
comm.Recv(Bdata, source=0, tag=2)

Cdata = numpy.add(Adata, Bdata)





comm.Send(Cdata, dest=0, tag=3)
```

```
status=MPI.status()  
data=comm.recv(source=MPI.ANY_SOURCE,  
               tag=MPI.ANY_TAG,  
               status=status)  
source=status.Get_source()  
tag=status.Get_tag()
```




WORKING EXAMPLES WITH MPI4PY

*Message Passing Interface
(Collective Communication)*

Broadcast**Scatter****Gather****All Gather**

Reduce



MPI_MAX – Returns the maximum element.

MPI_MIN – Returns the minimum element.

MPI_SUM – Sums the elements.

MPI_PROD – Multiplies all elements.

MPI LAND – Performs a logical “and” across the elements.



MPI_LOR – Performs a logical “or” across the elements.

MPI_MAXLOC –the maximum value and the rank of the process that owns it.

MPI_MINLOC –the minimum value and the rank of the process that owns it.

$$\pi = \int_0^1 \frac{4.0}{(1+x^2)} dx = \sum_{i=0}^N \frac{4.0}{(1+x^2)} \Delta x$$

$$\pi = \int_0^1 \frac{4.0}{(1+x^2)} dx = \sum_{i=0}^N \frac{4.0}{(1+x^2)} \Delta x$$



```
import time
def Pi(num_steps):
    start = time.time()
    step = 1.0/num_steps
    sum = 0
    for i in xrange(num_steps):
        x= (i+0.5)*step
        sum = sum + 4.0/(1.0+x*x)
    pi = step * sum
    end =time.time()
    print ("Pi with %d steps is %f in %f secs" %(num_steps, pi, end-start))

if __name__ == '__main__':
    Pi(100000000)
```

```
def Pi(num_steps):  
    comm = MPI.COMM_WORLD  
    rank = comm.Get_rank()  
    size = comm.Get_size()  
    Proc_num_steps = num_steps / size  
    Start = rank * Proc_num_steps  
    End = (rank + 1) * Proc_num_steps  
    print ("%d: [%d,%d]" % (rank, Start, End))  
  
    step = 1.0 / num_steps  
    sum = 0  
    for i in xrange(Start, End):  
        x = (i + 0.5) * step  
        sum = sum + 4.0 / (1.0 + x * x)  
    return sum
```

(...)

(...)

```
if __name__ == '__main__':  
    num_steps = 100000000  
    start = time.time()  
    localsum = Pi(num_steps)  
  
    localpi = localsum / num_steps  
    comm = MPI.COMM_WORLD  
    rank = comm.Get_rank()  
    pi = comm.reduce(localpi, op=MPI.SUM, root=0)  
    end = time.time()  
    if rank == 0:  
        print ("Pi with %d steps is %f in %f secs" % (num_steps, pi, end - start))
```



REFERENCES

ASTRONOMY DATA AND COMPUTING SERVICES

REFERENCES

- <https://computing.llnl.gov/tutorials/mpi/>
- <https://mpi4py.readthedocs.io/en/stable/>
- <https://cas-ereseach.github.io/astroinformatics2017/day3.html>

ADACS

ASTRONOMY DATA AND COMPUTING SERVICES